

COMP/ELEC 429


Introduction to Computer Networks

Lecture 10: Intra-domain routing

Slides from Edward W. Knightly, T. S. Eugene Ng, Ion Stoica, Hui Zhang

Viewing Routing as a Policy


- Given multiple alternative paths, how to route information to destinations should be viewed as a policy decision
- What are some possible policies?
 - Shortest path (RIP, OSPF)
 - Most load-balanced
 - QoS routing (satisfies app requirements)
 - etc



Internet Routing

- Internet topology roughly organized as a two level hierarchy
- First lower level – autonomous systems (AS's)
 - AS: region of network under a single administrative domain
- Each AS runs an intra-domain routing protocol
 - Distance Vector, e.g., Routing Information Protocol (RIP)
 - Link State, e.g., Open Shortest Path First (OSPF)
 - Possibly others
- Second level – inter-connected AS's
- Between AS's runs inter-domain routing protocols, e.g., Border Gateway Routing (BGP)
 - De facto standard today, BGP-4

Example




Intra-domain Routing Protocol

- Based on unreliable datagram delivery
- Link state
 - Open Shortest Path First (OSPF), based on Dijkstra's algorithm
 - Each router periodically floods *immediate* reachability information to other routers
 - Fast convergence, but high communication and computation overhead

Routing on a Graph


- Goal: determine a “good” path through the network from source to destination
 - Good often means the shortest path
- Network modeled as a graph
 - Routers ☾ nodes
 - Link ☾ edges
 - Edge cost: delay, congestion level,...



Link State Routing (OSPF): Flooding

- Each node knows its connectivity and cost to a direct neighbor
- Every node tells every other node this local connectivity/cost information
 - Via flooding
- In the end, every node learns the complete topology of the network
- E.g. A floods message


A connected to B cost 2
A connected to D cost 1
A connected to C cost 5




Flooding Details

- Each node periodically generates Link State Packet (LSP) contains
 - ID of node created LSP
 - List of direct neighbors and costs
 - Sequence number (64 bit, assume to never wrap around)
 - Time to live
- Flood is reliable
 - Use acknowledgement and retransmission
- Sequence number used to identify *newer* LSP
 - An older LSP is discarded
- Receiving node flood LSP to all its neighbors except the neighbor where the LSP came from
- LSP is also generated when a link's state changes (failed or restored)


Link State Flooding Example




Link State Flooding Example



Link State Flooding Example



Link State Flooding Example



A Link State Routing Algorithm

Dijkstra's algorithm

- Net topology, link costs known to all nodes
 - Accomplished via “link state flooding”
 - All nodes have **same** info
- Compute least cost paths from one node (“source”) to all other nodes
- Repeat for all sources

Notations

- $c(i,j)$: link cost from node i to j ; cost infinite if not direct neighbors
- $D(v)$: current value of cost of path from source to node v
- $p(v)$: predecessor node along path from source to v , that is next to v
- S : set of nodes whose least cost path definitively known

Dijkstra's Algorithm (A “Greedy” Algorithm)

1 **Initialization:**

2 $S = \{A\};$

3 for all nodes v

4 if v adjacent to A

5 then $D(v) = c(A,v);$

6 else $D(v) = \infty;$

7

8 **Loop**

9 find w not in S such that $D(w)$ is a minimum;

10 add w to S ;

11 update $D(v)$ for all v adjacent to w and not in S :

12 $D(v) = \min(D(v), D(w) + c(w,v));$


// new cost to v is either old cost to v or known

// shortest path cost to w plus cost from w to v

13 **until all nodes in S ;**

Example: Dijkstra's Algorithm

Step	start S	D(B),p(B)	D(C),p(C)	D(D),p(D)	D(E),p(E)	D(F),p(F)
→ 0	A	2,A	5,A	1,A	∞	∞
1						
2						
3						
4						
5						




```

1 Initialization:
2 S = {A};
3 for all nodes v
4   if v adjacent to A
5     then D(v) = c(A,v);
6     else D(v) =  $\infty$ ;
...
  
```

Example: Dijkstra's Algorithm

Step	start S	D(B),p(B)	D(C),p(C)	D(D),p(D)	D(E),p(E)	D(F),p(F)
0	A	2,A	5,A	1,A	∞	∞
→ 1	AD		4,D		2,D	∞
2						
3						
4						
5						




```

...
8  Loop
9  find w not in S s.t. D(w) is a minimum;
10 add w to S;
11 update D(v) for all v adjacent
   to w and not in S:
12   D(v) = min( D(v), D(w) + c(w,v) );
13 until all nodes in S;
  
```


Example: Dijkstra's Algorithm

Step	start S	D(B),p(B)	D(C),p(C)	D(D),p(D)	D(E),p(E)	D(F),p(F)
0	A	2,A	5,A	1,A	∞	∞
1	AD		4,D		2,D	∞
→ 2	ADE		3,E			4,E
3						
4						
5						




```

...
8  Loop
9   find w not in S s.t. D(w) is a minimum;
10  add w to S;
11  update D(v) for all v adjacent
    to w and not in S:
12    D(v) = min( D(v), D(w) + c(w,v) );
13  until all nodes in S;
    
```

Example: Dijkstra's Algorithm

Step	start S	D(B),p(B)	D(C),p(C)	D(D),p(D)	D(E),p(E)	D(F),p(F)
0	A	2,A	5,A	1,A	∞	∞
1	AD		4,D		2,D	∞
2	ADE		3,E			4,E
→ 3	ADEB					
4						
5						




```

...
8  Loop
9  find w not in S s.t. D(w) is a minimum;
10 add w to S;
11 update D(v) for all v adjacent
    to w and not in S:
12     D(v) = min( D(v), D(w) + c(w,v) );
13 until all nodes in S;
    
```

Example: Dijkstra's Algorithm

Step	start S	D(B),p(B)	D(C),p(C)	D(D),p(D)	D(E),p(E)	D(F),p(F)
0	A	2,A	5,A	1,A	∞	∞
1	AD		4,D		2,D	∞
2	ADE		3,E			4,E
3	ADEB					
→ 4	ADEBC					
5						




```

...
8  Loop
9  find w not in S s.t. D(w) is a minimum;
10 add w to S;
11 update D(v) for all v adjacent
    to w and not in S:
12     D(v) = min( D(v), D(w) + c(w,v) );
13 until all nodes in S;
    
```

Example: Dijkstra's Algorithm

Step	start S	D(B),p(B)	D(C),p(C)	D(D),p(D)	D(E),p(E)	D(F),p(F)
0	A	2,A	5,A	1,A	∞	∞
1	AD		4,D		2,D	∞
2	ADE		3,E			4,E
3	ADEB					
4	ADEBC					
→ 5	ADEBCF					



```

...
8  Loop
9   find w not in S s.t. D(w) is a minimum;
10  add w to S;
11  update D(v) for all v adjacent
    to w and not in S:
12     $D(v) = \min( D(v), D(w) + c(w,v) );$ 
13  until all nodes in S;
    
```

Overview of Assignment

Given a set of N routers, the goal is for EACH router to:

- (a) exchange HELLO packets with neighbours
- (b) create Link State Advertisement (LSA) packets based on neighboring nodes' info
- (c) broadcast the LSA packets to all other routers in the network
- (d) construct the network topology based on the LSA packets received from other routers,
- (e) determining the routing table entries based on this topology, by using Dijkstra' algorithm
 - (single source –all nodes shortest paths)
 - If multiple equal-cost paths exist, any one of them can be reported.

Overview of Assignment

- The HELLO packets will be exchanged every x seconds
- LSA updates will be sent every y seconds
- Routing table computation will be done every z seconds

`./ospf -i id -f infile -o outfile -h hi -a lsai -s spfi`

The values specified in the command line are:

- `-i id`: Node identifier value(i)
- `-f infile`: Input file
- `-o outfile`: Output file
- `-h hi`: HELLO INTERVAL (in seconds)
- `-a lsai`: LSA INTERVAL (in seconds)
- `-s spfi`: SPF INTERVAL (in seconds)

Overview of Assignment

5		10	
0	2	2	8
2	3	5	10
3	4	6	20
4	7	4	10
...			

The first entry on the first line specifies the number of routers (N) The node indices go from 0 to $(N - 1)$.

The second entry on the first line specifies the number of links.

Each subsequent row contains the tuple $(i, j, \text{MinC}_{ij}, \text{MaxC}_{ij})$. This implies a bidirectional link between nodes i and j .